KV Capital Sam Call Cheat Sheet

One-call, five-minute tactical notes for the AI Engineer hackathon problem-owner call. Use this live while dialing.

Call target	Sam - AI agent playing the problem owner. Phone: +1 587 816 8042. One call only. Five minutes maximum. Use the same phone number when submitting.
Submission constraint	Public GitHub repo + README + demo video <=3 minutes + phone number used for Sam call. Firm deadline: Friday, June 5, 2026 at 11:59 PM MST.

THE ONE OBJECTIVE

Extract the scoring rubric behind 'good': target user, final output, comp-selection priorities, data expectations, and scope boundaries. Do not over-talk. Do not sell too much. Ask, listen, recap.

Time	Action	Say / Do
0:00-0:20	Open	State your intended scope in one sentence.
0:20-4:20	Question sprint	Ask the six priority questions. Capture keywords only.
4:20-5:00	Recap	Repeat back the build direction and confirm no major miss.

Live opener - read this almost verbatim

- Hi Sam, this is Zhane. I am building the KV AI Engineer hackathon submission.

- I am planning to scope this as a residential comparable-sales workflow for Alberta home-builder financing.

- The workflow would input a subject property, generate/rank comparable sales, show explainable adjustments, and produce a short underwriting-ready comp summary.

- I have a few quick questions to make sure I am optimizing for the right workflow.

Your silent stance

- Deterministic first, AI-assisted second.

- Explainability beats hallucinated magic.

- A clean residential workflow beats a broad but shallow residential + commercial prototype.

- The final artifact should help an analyst or underwriter make a faster, auditable decision.

Your Positioning - Use Only If Helpful

Do not turn the call into a resume pitch. Use this only when Sam asks about approach, implementation, or why you are scoping the system this way.

15-SECOND SKILL PROOF

Compact version

- I build AI product systems and deterministic workflow tools as Technical Lead Developer at UMATTR, including Ralphplan AI.

- My strongest angle is full-stack AI workflow engineering: TypeScript/Next.js, Node, Supabase/Postgres, Vercel, RAG/RAI, eval/QA loops, observability, and agentic workflows.

- Because I also build Solidity/state-machine systems, I think in audit trails, source-of-truth events, deterministic rules, and explainable decision paths.

HOW YOUR SKILLS MAP TO THEIR PROBLEM

Your strength	How to connect it to comp analysis
Agentic AI workflows	Use AI for candidate retrieval, explanation, memo drafting, and analyst-assist workflows - not as an unchecked valuation oracle.
RAG/RAI + context construction	Ground outputs in comparable records, listing fields, methodology notes, and explicit assumptions.
Eval/QA loops	Test comp-ranking outputs against expected cases: near match, bad location, stale sale, size mismatch, property-type mismatch.
Observability	Log every workflow step: property loaded, comps generated, scores computed, flags raised, memo generated.
Full-stack product execution	Build a clean deployed Vercel app with a usable UI, README, demo, and credible architecture.
Solidity/Web3 proof instincts	Bring auditability: immutable-style event trail, confidence bands, evidence references, deterministic scoring rules.
Uranium/nuclear workflow modeling	Frame the system like industrial decision support: structured inputs, constraints, margins, risk flags, human review.

LANGUAGE GUARDRAILS

Avoid	Say instead
I will fully automate underwriting.	I am building an analyst-assist workflow that keeps human review in the loop.
The AI will determine true property value.	The system ranks comps deterministically and drafts reviewable explanations.
I can scrape everything live in a weekend.	The demo uses synthetic/public-like Alberta data with clear methodology and limitations.
I will build a huge residential + commercial platform by Friday.	I will ship a focused residential MVP, with commercial as a future extension.

Six Priority Questions

Ask these in order. If Sam gives a long answer, skip lower-priority questions and preserve time for the recap.

#	Question	Why it matters
1	For this hiring exercise, what does a strong submission demonstrate most: valuation accuracy, workflow design, AI reasoning, source traceability, UI quality, or engineering completeness?	This gives the rubric. Build toward this answer.
2	Who should I design the workflow for: an underwriter, analyst, loan officer, or someone else reviewing a builder financing request?	This sets labels, UX, and final output language.
3	What is the ideal final output of comp analysis: a ranked comp table, estimated value range, risk flags, written memo, or something else?	This defines the final screen and README story.
4	When choosing comparable properties, what matching factors matter most: location, property type, square footage, bedrooms, age, sale date, price per square foot, or builder/project context?	This sets scoring weights.
5	Is it better to use synthetic/public-like data with a clear methodology, or prioritize finding a real public Alberta dataset even if limited?	This prevents wasting build time on fragile data gathering.
6	For a solo build, would you rather see a tight residential workflow polished end-to-end, or a broader prototype that also sketches commercial borrower support?	This protects against overbuilding.

FALLBACK PROBES

Use only if an answer is vague

- When an analyst rejects a comp, what is the usual reason?

- What makes a comp dangerous or misleading?

- What part of the current workflow consumes the most human time?

- Where does human judgment still need to stay in the loop?

- Should the prototype optimize for speed, trust, or final valuation precision?

CLOSING RECAP

Read this in the final 40 seconds

- Perfect. I will scope this as a focused residential comp-analysis prototype.

- The core flow will be: subject property input, comparable candidate generation, ranked comps, explainable scoring, risk flags, value range, and an underwriting-style summary with an audit trail.

- I will keep the README explicit about assumptions, limitations, and what I would build next.

Product Hypothesis To Validate

This is your default plan unless Sam changes the priority order.

WORKING PRODUCT NAME

Recommended	BuilderComp AI - AI-assisted residential comp-analysis workflow for Alberta builder financing.

CORE WORKFLOW

Step	What user sees	What system does
1. Subject property	Enter/view address, beds, baths, sqft, type, year built, asking price/estimated value.	Normalize fields and create comparison profile.
2. Candidate comps	Show 6-10 recent comparable properties.	Filter by location, type, size, sale recency, and core attributes.
3. Deterministic ranking	Score each comp and show top matches.	Compute weighted similarity score and apply risk penalties.
4. Explainability	Why this comp was selected or discounted.	Generate readable reasoning from deterministic scoring facts.
5. Risk flags	Warnings: stale sale, distance, size mismatch, property-type mismatch, outlier price.	Surface reasons for human review.
6. Value range	Low/base/high range with confidence.	Estimate value from adjusted comp signals.
7. Memo output	Underwriting-ready summary.	Create reviewable narrative with evidence and limitations.
8. Event trail	Live workflow log.	Record auditable steps and decisions.

WHAT TO BUILD IF SAM PRIORITIZES...

Sam says	Your build response
Explainability	Emphasize comp reasoning panel, flags, weight breakdown, and memo traceability.
Speed	Emphasize one-screen workflow, fast candidate ranking, and automatic memo generation.
Accuracy	Emphasize scoring methodology, weight tuning, outlier handling, and known limitations.
Data realism	Use better public-like synthetic dataset and cite methodology; avoid pretending it is MLS-grade.
Commercial extension	Keep residential as MVP; add one commercial example as future/bonus tab only after core works.

Deterministic System Notes

This is the engineering backbone. If Sam asks what you mean by AI, explain that the ranking engine should be auditable first, then AI can narrate and accelerate the workflow.

DRAFT SCORING MODEL

Factor	Draft weight	Reason	Question for Sam
Location / distance	25%	Nearby sales usually carry strong signal.	What distance radius is reasonable for Alberta residential builder lending?
Property type	15%	Detached, duplex, townhouse, condo must not mix blindly.	Which type mismatch should disqualify a comp?
Size / sqft	15%	Major size gaps distort price comparison.	What sqft tolerance is acceptable?
Beds / baths	10%	Layout similarity matters for residential buyer behavior.	Should beds/baths be hard filters or soft penalties?
Sale recency	15%	Stale sales need time adjustment or lower confidence.	What sale window is acceptable? 3, 6, 12 months?
Price per sqft	10%	Useful for outlier detection and normalization.	Should PPSF outliers be excluded or flagged?
Age / build quality	5%	Newer homes differ in finishes, warranties, and buyer premium.	Does year built matter more than condition/finish?
Risk flags	5%	Keep humans aware of weak assumptions.	Which flags matter most to KV?

NON-NEGOTIABLE ARCHITECTURE PRINCIPLE

Deterministic first

- The score should come from visible inputs and rules, not a hidden LLM opinion.

- AI can summarize why the score happened and draft the memo.

- Every comp should have a reason, adjustment, confidence, and flag state.

- Every major step should be event-logged for auditability.

LIKELY README ARCHITECTURE SENTENCE

BuilderComp AI separates deterministic comp ranking from AI-assisted explanation. The scoring engine ranks candidate properties from structured inputs, while the reasoning layer turns those facts into an analyst-readable summary, risk review, and underwriting memo.

Call Capture Sheet

Write short keywords while listening. Do not try to transcribe everything.

Decision area	Sam's answer / keywords
Success criteria	
Target user	
Ideal final output	
Most important comp factors	
Data preference	
Scope: residential vs commercial	

IMMEDIATELY AFTER CALL

Convert transcript into build decisions

- Rewrite README problem understanding in Sam's language.
- Lock scoring weights or explain why they are draft defaults.
- Update demo script to mirror the final output Sam described.
- Add one screenshot or UI section that directly proves the main success criterion.
- Do not add new big features unless Sam clearly identifies a critical missing workflow step.

Post-Call Build Lock

Use this page to prevent scope creep after the call.

DEFAULT MVP AFTER CALL

Ship these before polish

- Subject property input/view
- Comparable candidates list
- Deterministic scoring + weight breakdown
- Risk flags and outlier detection
- AI-style explanation panel grounded in scores
- Value range estimate
- Underwriting memo output
- Workflow Events Tracker / audit trail
- README with assumptions, how to run, architecture, limitations, next steps
- Vercel deploy + public GitHub + <=3 minute demo video

NO-GO LIST

Avoid after the call unless directly required

- Dragging map system
- Real MLS scraping
- Auth/payment/database complexity
- Huge commercial workflow
- Unverified live data claims
- A generic chatbot that cannot show why a comp was selected
- Major UI rewrite after Wednesday night

FINAL BUILD THESIS

A credible hiring submission is a small, complete, auditable product: not an overbuilt app, not a fake valuation oracle, and not a flashy UI with no methodology. Your strongest angle is a deterministic AI workflow console that helps KV analysts move from subject property to comp packet faster, with explainable ranking, risk flags, and a reviewable memo.

Internal deadline	MVP locked Wednesday night. Polish/demo Thursday. Submit Friday by 8:00 PM Edmonton time. Hard external deadline: Friday, June 5, 2026 at 11:59 PM MST.